

1. Introduction to GANs

- Generative Adversarial Networks (GANs) are a class of machine learning models introduced by Ian Goodfellow in 2014. GANs consist of two neural networks that are trained simultaneously through adversarial learning:
 - **Generator:** A model that generates synthetic data resembling the real data.
 - **Discriminator:** A model that distinguishes between real and fake data.
- The Generator learns to produce outputs that can "fool" the Discriminator, while the Discriminator improves its ability to classify real versus fake data. This adversarial process helps the Generator create highly realistic data over time.
- Applications of GANs include image synthesis, data augmentation, video generation, and style transfer.

2. Dataset Description and Preprocessing Steps

→ Dataset:

The MNIST dataset was used for this project. It consists of:

- **Training Data:** 60,000 images of handwritten digits.
- **Test Data:** 10,000 images of handwritten digits.
- **Image Dimensions:** 28x28 pixels, grayscale.
- **Classes:** 10 (digits 0–9).

→ Preprocessing Steps:

1. Loading the Dataset:

- The dataset was loaded using TensorFlow's built-in `tf.keras.datasets.mnist` function.

2. Normalization:

- Pixel values were scaled to the range `[-1, 1]` for faster and more stable GAN training.

3. Batching and Shuffling:

- Data was reshaped to include a channel dimension and batched into groups of 256 images for efficient training.

```
# Load and preprocess the MNIST dataset
(train_images, _), (_, _) = tf.keras.datasets.mnist.load_data()

# Normalize the images to the range [-1, 1]
train_images = train_images.astype("float32") / 255.0
train_images = (train_images - 0.5) / 0.5

# Reshape the images to include the channel dimension (28x28x1)
train_images = np.expand_dims(train_images, axis=-1)
```

✓ 0.4s

3. Detailed Explanation of Model Architecture

→ Generator:

The Generator maps a random noise vector to a synthetic image resembling the MNIST digits. The architecture includes:

- **Dense Layer:** Projects the input noise vector to a 7x7x256 feature map.
- **Batch Normalization:** Stabilizes training by normalizing activations.
- **Transposed Convolution Layers:** Upsample the feature map to 28x28.
- **Activation Functions:** LeakyReLU for hidden layers, tanh for the output layer (to match normalized image range).

```
generator.summary()
```

✓ 0.0s

Model: "sequential_9"

Layer (type)	Output Shape	Param #
=====		
dense_9 (Dense)	(None, 12544)	1254400
batch_normalization_15 (Batch Normalization)	(None, 12544)	50176
leaky_re_lu_23 (LeakyReLU)	(None, 12544)	0
reshape_5 (Reshape)	(None, 7, 7, 256)	0
conv2d_transpose_15 (Conv2DTranspose)	(None, 7, 7, 128)	819200
batch_normalization_16 (Batch Normalization)	(None, 7, 7, 128)	512
leaky_re_lu_24 (LeakyReLU)	(None, 7, 7, 128)	0
conv2d_transpose_16 (Conv2DTranspose)	(None, 14, 14, 64)	204800
batch_normalization_17 (Batch Normalization)	(None, 14, 14, 64)	256
...		
Total params: 2330944 (8.89 MB)		
Trainable params: 2305472 (8.79 MB)		
Non-trainable params: 25472 (99.50 KB)		

→ Discriminator:

- The Discriminator evaluates whether an input image is real or fake. The architecture includes:
 - **Convolutional Layers:** Downsample the input image.
 - **Dropout:** Prevents overfitting by randomly deactivating neurons during training.
 - **Dense Output Layer:** Produces a single value indicating real or fake classification.

```
discriminator.summary()
```

✓ 0.0s

Model: "sequential_10"

Layer (type)	Output Shape	Param #
=====	=====	=====
gaussian_noise_2 (Gaussian Noise)	(1, 28, 28, 1)	0
conv2d_8 (Conv2D)	(1, 14, 14, 64)	1664
leaky_re_lu_26 (LeakyReLU)	(1, 14, 14, 64)	0
dropout_8 (Dropout)	(1, 14, 14, 64)	0
conv2d_9 (Conv2D)	(1, 7, 7, 128)	204928
leaky_re_lu_27 (LeakyReLU)	(1, 7, 7, 128)	0
dropout_9 (Dropout)	(1, 7, 7, 128)	0
flatten_4 (Flatten)	(1, 6272)	0
dense_10 (Dense)	(1, 1)	6273

=====

Total params: 212865 (831.50 KB)

Trainable params: 212865 (831.50 KB)

Non-trainable params: 0 (0.00 Byte)

4. Explanation of Loss Functions and Training Strategy

→ Loss Functions:

1. Discriminator Loss:

- Measures the Discriminator's ability to classify real and fake images correctly.
- Combines loss for real and fake images.

2. Generator Loss:

- Measures the Generator's ability to fool the Discriminator.

→ Training Strategy

- **Adversarial Training:**
 - Alternate updates between the Generator and Discriminator.
 - Use of Adam optimizer with learning rates of $1e-4$ for both models.
- **Training Duration:**
 - The models were trained for 100 epochs, with Generator outputs visualized after each epoch.


5. Evaluation and Visualization of Results

→ Training Progress:

- The Generator's outputs were visualized at various epochs to observe learning progression. Early outputs were noisy and lacked structure, but later outputs resembled handwritten digits.

→ Loss Graph:


- The loss trends for the Generator and Discriminator are shown below:




→ Generated Image Results

- Below are sample images generated by the GAN over 100 epochs improving:


Generated Images at Epoch 1



Generated Images at Epoch 25



Generated Images at Epoch 50



Generated Images at Epoch 75



Generated Images at Epoch 100



6. Conclusion

- The GAN successfully learned to generate synthetic handwritten digits that resemble the MNIST dataset. Key observations include:
 - **Generator Improvements:** Over time, the Generator produced clearer and more realistic digits.
 - **Adversarial Learning:** The balance between Generator and Discriminator losses indicates stable training.
- **Future Improvements:**
 - Train for more epochs for further refinement.
 - Use Conditional GANs to generate specific digits.
 - Experiment with advanced architectures like DCGANs or StyleGANs.
- This project demonstrates the fundamentals of GANs and their potential for generating realistic data.